

Tutorial 3


COMP4621

Concurrent Client Server Using fork

Spring 2024



TA: Xinyu YANG
Email: xinyu.yang@connect.ust.hk




Outline

- Understanding the multiple TCP client server problem
- Review fork() function
- Example of concurrent TCP client server problem
- Set timeout using setsockopt() function

Understanding the multiple client server problem

- The figure shows several clients, which have contacted one server. The client connections conceptually form spokes around the central server.



Ways to achieve concurrent client server

- Forked server processes (multi-process method)
- Threaded server processes (multi-thread method)
- One process and a `select(2)` call
- One process and a `poll(2)` call

Comparison of the four ways

Method	Advantages	Disadvantages
Forked server process	<ul style="list-style-type: none">• Simplest	<ul style="list-style-type: none">• Sharing information becomes more complex.• Requires more CPU to start and manage a new process for each request
Threaded server process	<ul style="list-style-type: none">• Lightweight advantages of the multi-process method	<ul style="list-style-type: none">• Difficult to debug

- Select(2) and poll(2) system calls offer a different way to block execution of the server until an event occurs (usually receiving data on a socket)

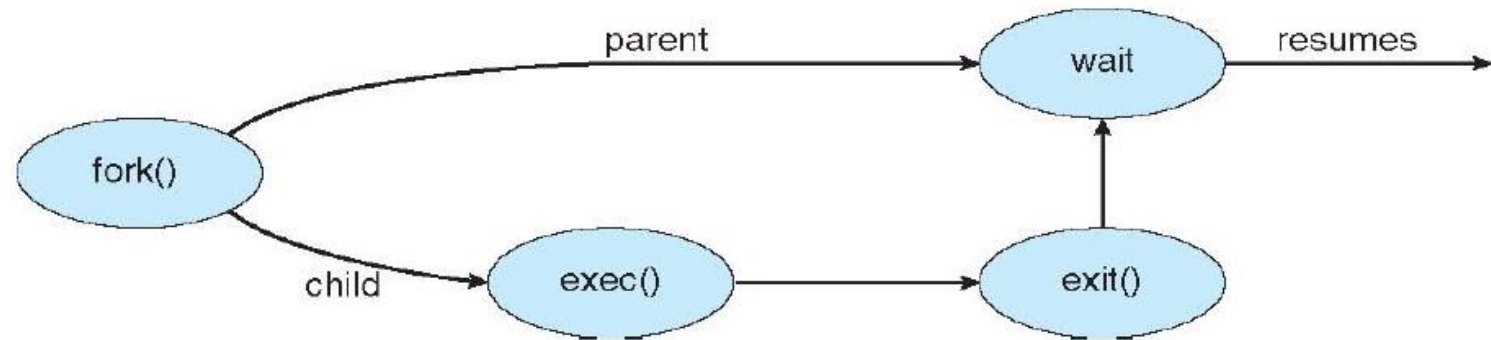
Review of fork()

- fork – create a child process
- SYNOPSIS

```
#include <sys/types.h>
#include <unistd.h>


pid_t fork(void)
```

- DESCRIPTION:
 - fork() creates a new process by duplicating the calling process. The new process is referred to as the child process. The calling process is referred to as the parent process.
 - The child process and the parent process run in separate address spaces (i.e, they do not share their data).




Return values of fork

■ Successful



■ Not Successful



errno is set to indicate error

fork example

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <unistd.h>
#include <sys/types.h>

#define BUFSIZE 1024
int main(int argc, char *argv[])
{
    char buf[BUFSIZE];
    size_t readlen, writelen, slen;
    pid_t cpid, mypid;
    pid_t pid = getpid();          /* get current processes PID */
    printf("Parent pid: %d\n", pid);
    cpid = fork();
    if (cpid > 0) {                /* Parent Process */
        mypid = getpid();
        printf("[%d] parent of [%d]\n", mypid, cpid);
    } else if (cpid == 0) {        /* Child Process */
        mypid = getpid();
        printf("[%d] child\n", mypid);
    } else {
        perror("Fork failed");
        exit(1);
    }
    exit(0);
}
```


fork example

Parent Process
(cpid=childpid)

```
cpid = fork();
if (cpid > 0) {      /* Parent Process */
    mypid = getpid();
    printf("[%d] parent of [%d]\n", mypid, cpid);
} else if (cpid == 0) { /* Child Process */
    mypid = getpid();
    printf("[%d] child\n", mypid);
} else {
    perror("Fork failed");
    exit(1);
}
```

Child Process
(cpid=0)

```
cpid = fork();
if (cpid > 0) {      /* Parent Process */
    mypid = getpid();
    printf("[%d] parent of [%d]\n", mypid, cpid);
} else if (cpid == 0) { /* Child Process */
    mypid = getpid();
    printf("[%d] child\n", mypid);
} else {
    perror("Fork failed");
    exit(1);
}
```

```
cpid = fork();
if (cpid > 0) {      /* Parent Process */
    mypid = getpid();
    printf("[%d] parent of [%d]\n", mypid, cpid);
} else if (cpid == 0) { /* Child Process */
    mypid = getpid();
    printf("[%d] child\n", mypid);
} else {
    perror("Fork failed");
    exit(1);
}
```

wait() and waitpid()

- **pid_t wait(int *status);**
- **pid_t waitpid(pid_t pid, int *status, int options);**
- Description:
 - The **wait()** system call suspends execution of the calling process until one of its children terminates
 - The **waitpid()** system call suspends execution of the calling process until a child specified by *pid* argument has changed state.
- Arguments:
 - PID can be <-1 meaning wait for any child process whose process group ID is equal to the absolute value of *pid*; -1 meaning wait for any child process; 0 meaning wait for any child process whose process group ID is equal to that of the calling process; >0 meaning wait for the child whose process ID is equal to the value of *pid*.
 - OPTIONS can be 0, WNOHANG, ... refer to <https://linux.die.net/man/2/waitpid>.
- Return value:
 - **wait()**: on success, returns the process ID of the terminated child; on error, -1 is returned.
 - **waitpid()**: on success, returns the process ID of the child whose state has changed; if WNOHANG was specified and one or more child(ren) specified by *pid* exist, but have not yet changed state, then 0 is returned. On error, -1 is returned.

Example of multiple client server

- Clients:

- User types line of text

- Client program sends line to server

- Server:

- Server receives line of text

- Capitalizes all the letters

- Sends modified line to client

- Clients:

- Receives line of text

- Displays

Skeleton of multiple client server using fork()

```
// Create a TCP/IP socket
listener = socket(PF_INET, SOCK_STREAM, 0);
// Bind the server address:
res = bind(listener, (struct sockaddr *)&server_addr, sizeof(server_addr));
// Make a listening socket:
res = listen(s, 10);
```

Skeleton of multiple client server using fork()

```
// Start the server loop
for (;;) {
    // Wait for a connect
    client_socket = accept(listener, (struct sockaddr *)&client_addr, sizeof(client_addr));
    // Fork a new server process to service this client
    if ((PID = fork()) == -1) {
        // Failed to fork: Give up
        close(client_socket);
        continue;
    } else if (PID > 0) { /* The parent process */
        // Parent process ... so does not need the client socket anymore
        close(client_socket);
        continue;
    }
    // Child process we can close the server socket s
    close(listener); /* The child process */
    /*Exchange data with the Client*/
    // TODO
    // Close the client socket
    close(client_socket);
    // Child process must exit:
    exit(0);
}
```

Skeleton of multiple client server using fork()

```
static void
sigchld_handler(int signo){
    pid_t PID;
    int status;

    do {
        PID = waitpid(-1, &status, WNOHANG);
    } while (PID != -1);

    /* Re-instate handler*/
    signal(SIGCHLD, sigchld_handler);
}
```

- Understanding Process Termination Processing
- 1. The signal SIGCHLD is sent by the kernel to the parent upon termination of its child process
- Normally SIGCHLD signal is ignored by default
- 2. But we can register a signal handler to process it: sigchld_handler() is called, because the function was registered for the SIGCHLD signal (“signal(SIGCHLD, sigchld_handler)”)
- 3. The sigchld_handler() executes a loop calling waitpid(2) until no more exit status information is available.
- 4. The SIGCHLD handler is re-instated

There's a problem herein. Find and modify it in your code.

Test the result

```
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$ ./cc_server  
Start communicating with client, IP address is: 127.0.0.1, port is: 35708  
Start communicating with client, IP address is: 127.0.0.1, port is: 48742  
From client: asdf  
To client: ASDF  
From client: lkjh  
To client: LKJH
```

```
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$ ./cc_client  
Socket successfully created...  
Connected to the server...  
Enter the string: asdf  
From Server : ASDF  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$
```

```
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$ ./cc_client  
Socket successfully created...  
Connected to the server...  
Enter the string: lkjh  
From Server : LKJH  
[xyangcp@csl2wk31 ~/comp4621/tutorial3]$
```

Non-blocking sockets with timed I/O

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int setsockopt(int s, int level, int optname, const void *optval, socklen_t optlen);
```

- All the socket information is in the man page for socket in section 7. (Type: “man 7 socket” to get all the information)
- Arguments:
 - s is the socket you’re talking about
 - level should be set to SOL_SOCKET
 - optname is the name of the option you’re interested in changing
 - optval is usually a pointer to an int indicating the value in question. For booleans, zero is false, and non-zero is true
 - optlen should be set to the length of optval, probably sizeof(int), but varies depending on the option.
- Return values
 - Returns zero on success, or -1 on error (and errno will be set accordingly)

Non-blocking sockets with timed I/O

- Socket options
 - `SO_RCVTIMEO` and `SO_SNDTIMEO`
 - specify the receiving or sending timeouts until reporting an error.
 - The argument is a struct `timeval`.
 - If an input or output function blocks for this period of time, and data has been sent or received, the return value of that function will be the amount of data transferred; if no data has been transferred and the timeout has been reached, then -1 is returned with `errno` set to `EAGAIN` or `EWOULDBLOCK`, or `EINPROGRESS` just as if the socket was specified to be nonblocking. If the timeout is set to zero (the default), then the operation will never timeout.

```
struct timeval {  
    int tv_sec;      // seconds  
    int tv_usec;    // microseconds  
};
```

Non-blocking sockets with timed I/O

What's the problem in non-blocking I/O?

- Example

```
tv.tv_sec = 5; /* 5 Secs Timeout */
tv.tv_usec = 0;
ret = setsockopt(client_socket, SOL_SOCKET, SO_RCVTIMEO, (struct timeval *)&tv, sizeof(struct timeval));

if (ret == SO_ERROR)
{
    printf("setsockopt() failed...\n");
    return -1;
}
else
    printf("setsockopt() is OK!\n");

// read the message from client and copy it in buffer
int numBytes = recv(client_socket, buffer, sizeof(buffer), 0);
if (numBytes <= 0)
{
    // nothing received from client in last 5 seconds
    close(client_socket);
    printf("nothing received in the last %ld seconds, close client socket\n", tv.tv_sec);
    break;
}
```

Summary & Tasks

Summary

- We talked about how to implement a concurrent server using `fork()`.
- We learned how to set timeout for a socket using `setsockopt()`.

Task

The code for concurrent server and client are available in Canvas. Run the code and answer the following questions:

1. Can the UDP server in Tutorial2 serves multiple UDP clients concurrently? Why?
2. What will happen if we do not close the `server_socket/client_socket` in client/server process?
3. What's the problem in non-blocking I/O?
4. Find the problem of `sigchld_handler` function and give the right one.

Submit a pdf document on Canvas before 11:59pm tonight!

Thank You!

